

Large scale deep learning in PFN: from 15-min Imagenet to PFDet

Hirochika Asai, Keisuke Fukuda

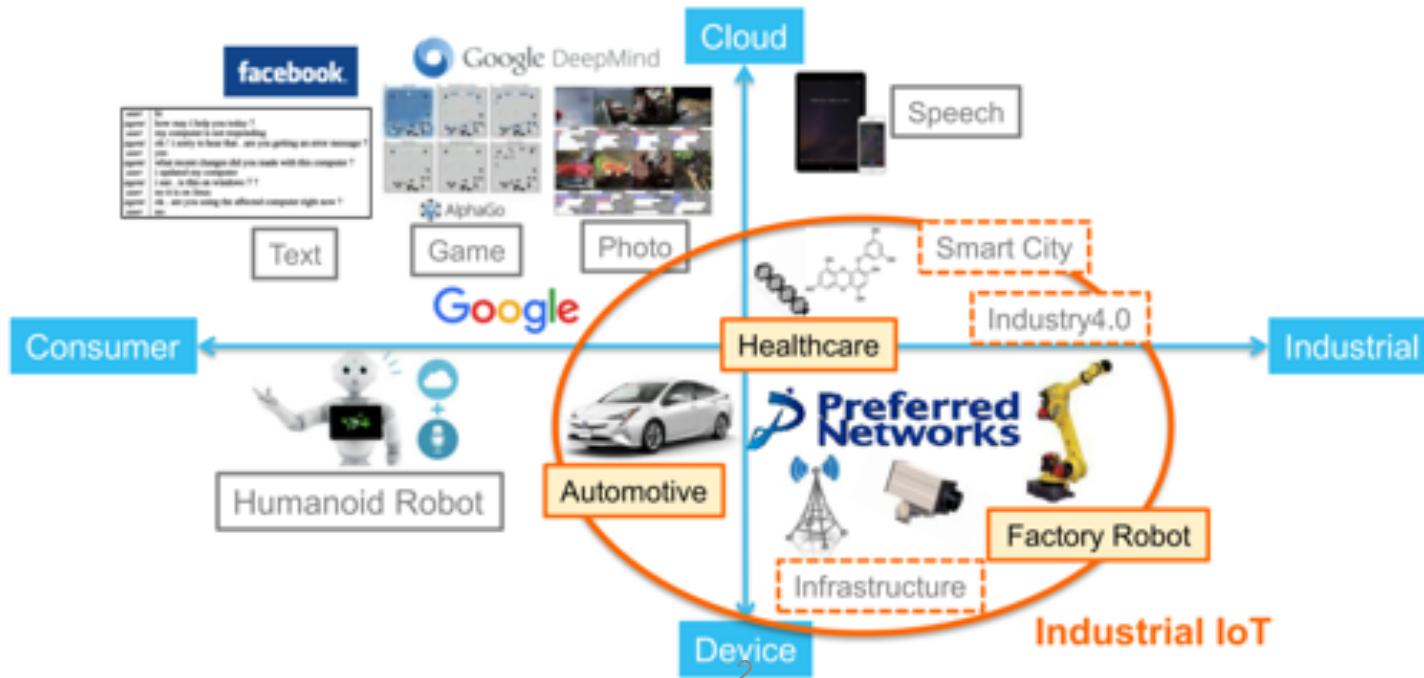
Preferred Networks, Inc.



Who we are?

Preferred Networks, Inc. (PFN):

A Tokyo-based Deep Learning & IoT company



Our Strategic Partners



Hakuhodo DY holdings



Roche A member of the Roche group

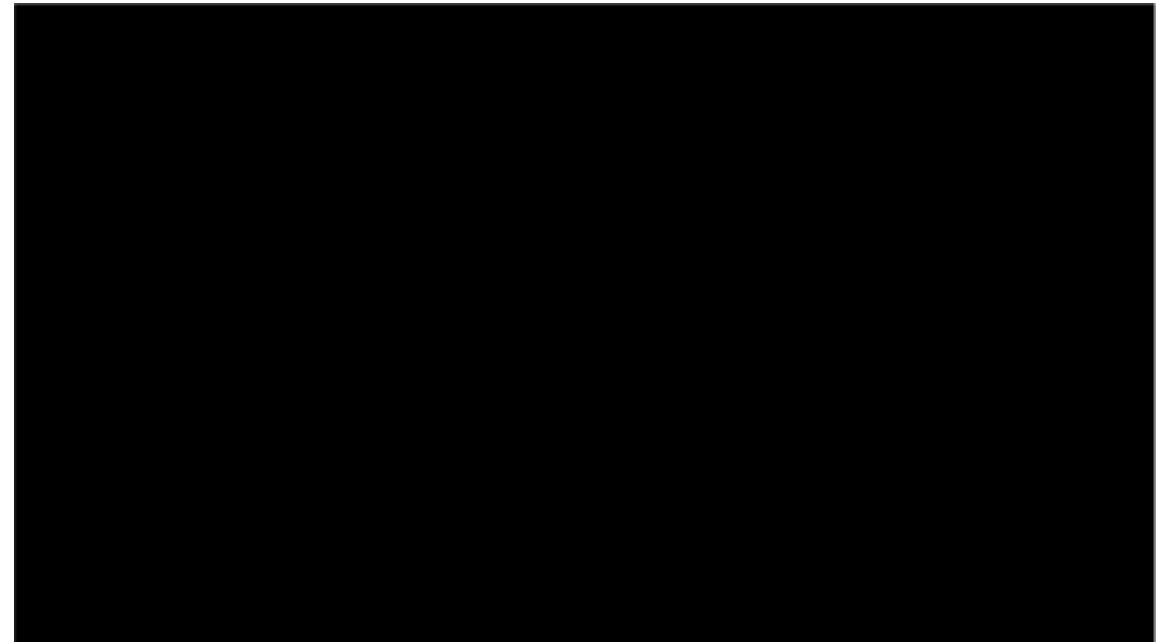


and Collaborators



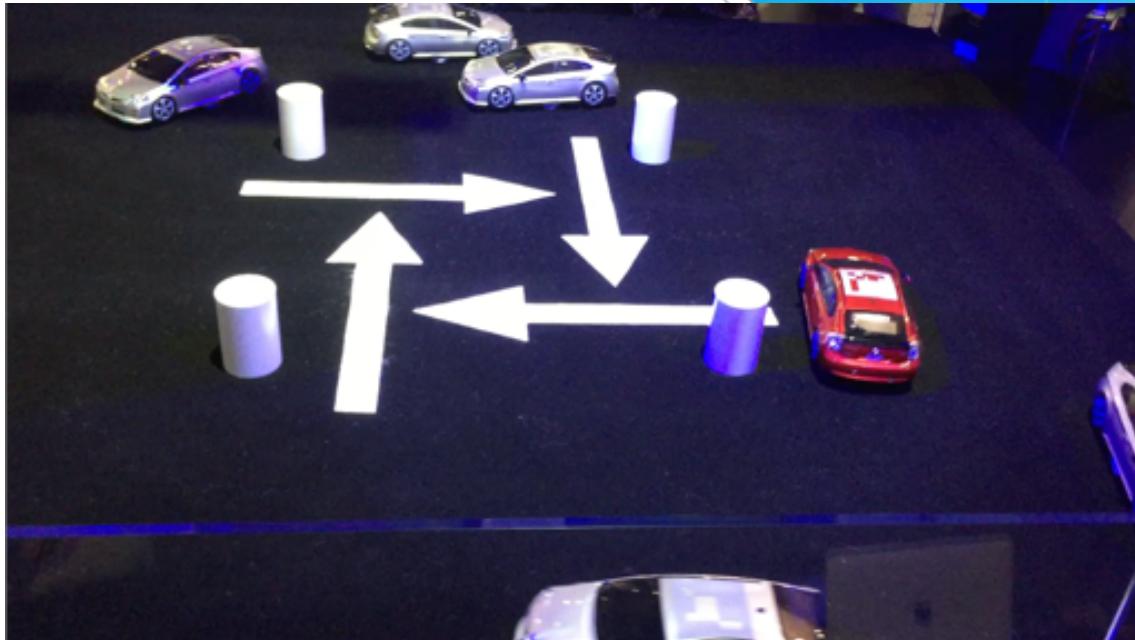
2015: OPTIMIZATION OF BIN-PICKING FANUC ROBOTS

- Picking random object is a typical task that is “easy for human, hard for robots”.



@CES 2016: CARS THAT DON'T CRASH

- Car positions are tracked from a ceiling camera and each car is controlled individually.
- White cars are autonomous
- The red car is a manually-controlled “evil” car: trying to disrupt other cars



@ICRA 2017 VOICE RECOGNITION + OBJECT PICKING

“Interactively Picking Real-World Objects with Unconstrained Spoken Language Instructions”
[arXiv:1710.06280](https://arxiv.org/abs/1710.06280)

- ICRA is a top-tier conference on robotics
- Best Paper Award on Human-Robot Interaction
- Technologies:
 - Visual recognition
 - Natural language processing (NLP)
- The robot can understand ambiguous words:
 - “The Teddy bear”
 - “The brown fluffy stuff”



<https://projects.preferred.jp/tidying-up-robot/>

@CEATEC JAPAN 2018 Autonomous Tidying-up Robot System



Technologies behind the demos: Distributed Deep Learning

MN-1: An in-house supercomputer

- **MN-1a** (Sep. '17~)
 - 1024 NVIDIA Tesla P100
 - InfiniBand FDR
 - Peak 9.3 Peta FLOPS (SP)
 - #227 in Top500 Nov. 2018
- **MN-1b** (July. '18~)
 - 512 NVIDIA Tesla V100 32GB
 - InfiniBand EDR
 - Peak 56 Peta (tensor) Flops
- Targeting Exa FL ops by 2020





Chainer

A Powerful, Flexible, and Intuitive Framework for Neural Networks

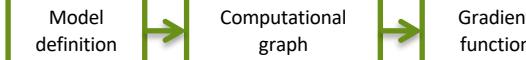
[GET STARTED](#)[LEARN MORE](#)

<https://chainer.org/>

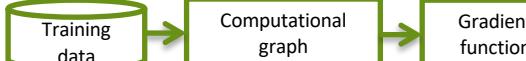
Chainer: A Flexible Deep Learning Framework

Define-and-Run

Define



Run

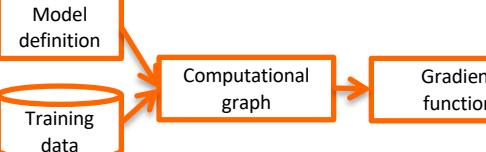


Caffe2, TensorFlow etc.

Define-by-Run



Define-by-Run



ChainerMN: Distributed Training with Chainer

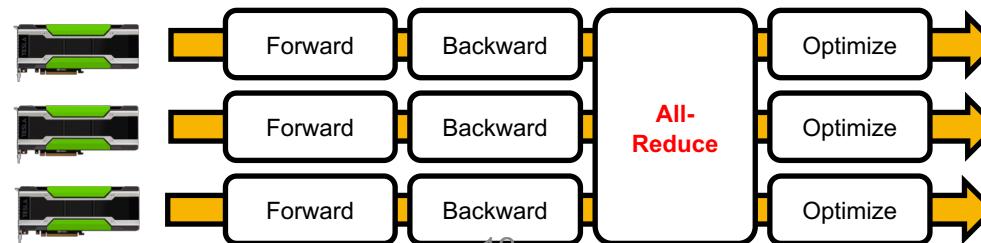
- Add-on package for Chainer
- Enables multi-node distributed deep learning using NVIDIA NCCL2

Features

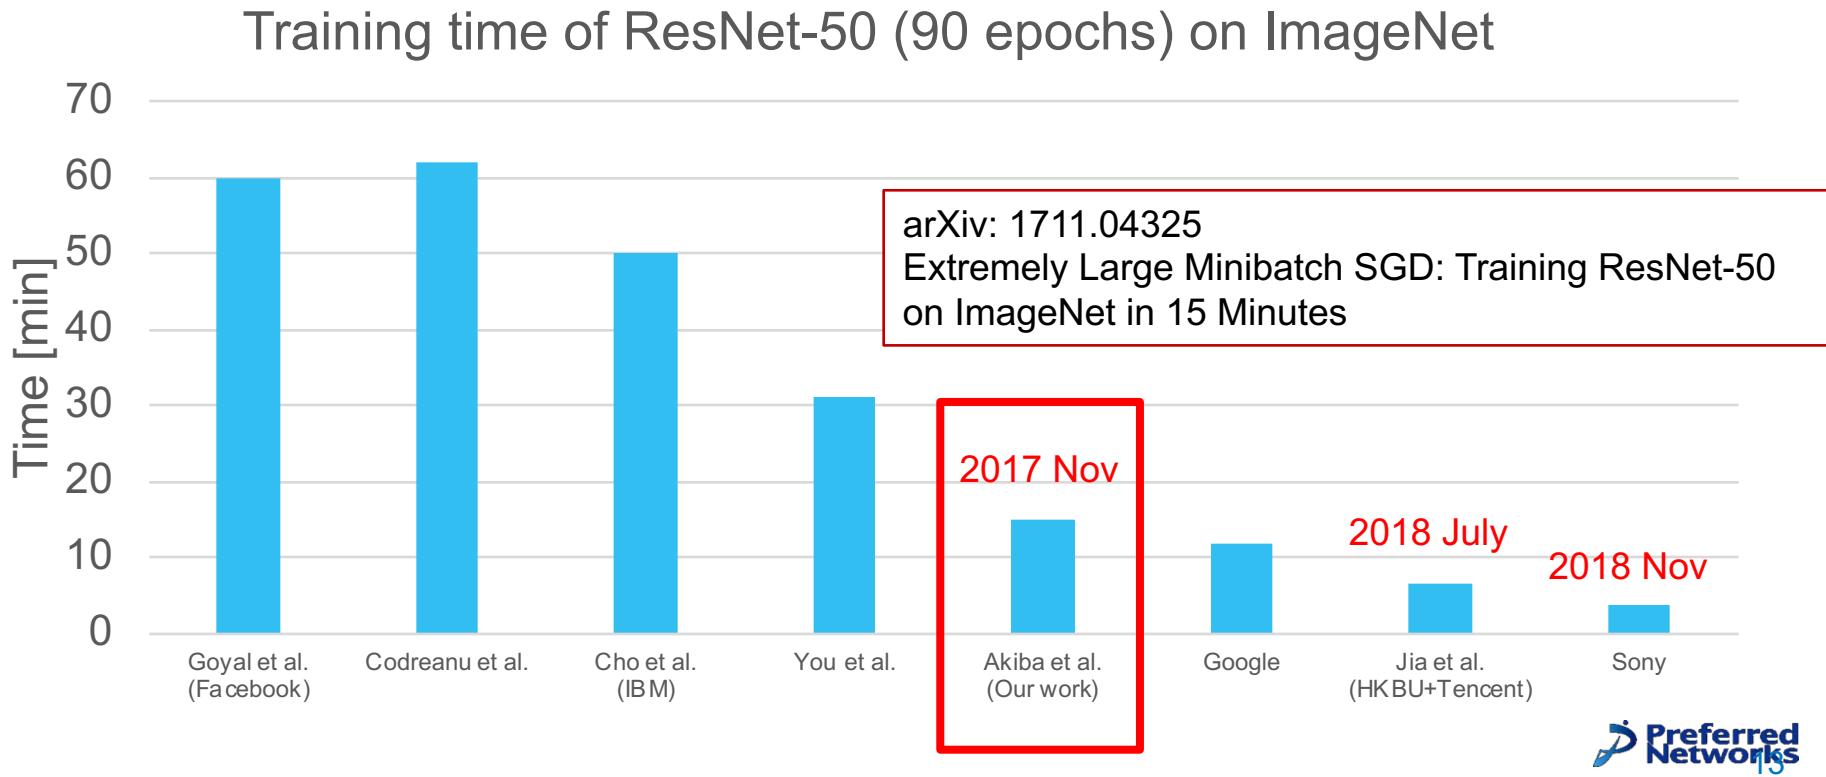
- **Scalable:** Near-linear scaling with hundreds of GPUs
- **Flexible:** Even GANs, dynamic NNs, and RL are applicable



Distributed Training with ChainerMN



Achievement on MN-1a: ImageNet in 15 minutes



NVIDIA

NVIDIA POWERS WORLD'S FASTEST
DEEP LEARNING PERFORMANCE



Jen-Hsun Huang
NVIDIA CEO, at SC'17

Achievement on MN-1b: PFDet in OIC 2018

The screenshot shows the competition page for the Google AI Open Images - Object Detection Track. At the top, there's a banner with the text "Featured Prediction Competition" and "Google AI Open Images - Object Detection Track". Below the banner, it says "Detect objects in varied and complex images." On the right side of the banner, it shows "\$30,000 Prize Money". Below the banner, there's a logo for Google AI and the text "454 teams · 2 months ago". At the bottom of the main section, there are navigation links: Overview (which is underlined), Data, Kernels, Discussion, Leaderboard, and Rules.

Overview

Description

Introduction

Google AI (Google's AI research arm, tasked with advancing AI for everyone) is challenging you to build an algorithm that detects objects automatically using an absolutely massive training dataset — one with more varied and complex bounding-box annotations and object classes than ever before.

Evaluation

Prizes

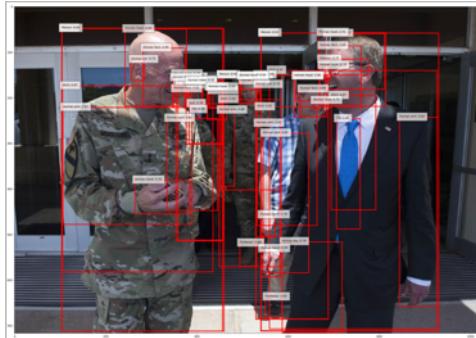
Timeline

Here's the background. Computers are getting better and better at vision. But in a few critical ways, they still can't match a human's intuitive perception.

For example, what do you see when you look at this photo?

Achievement on **MN-1b**: PFDet in OIC 2018

- Google AI Open Images - Object Detection Track
 - Competition using Largest-class image dataset
 - 12 million bounding boxes, 1.7 million images
 - 454 competitors
 - Approx. 500GB (annotated subset)
- Object detection: much harder than object recognition task

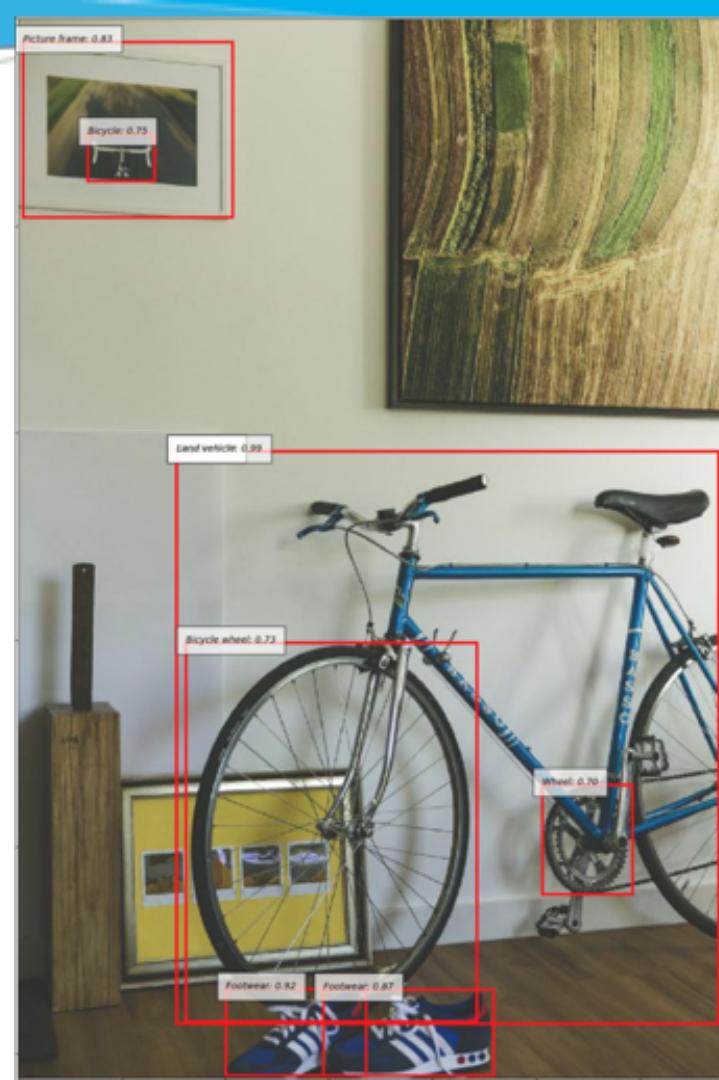
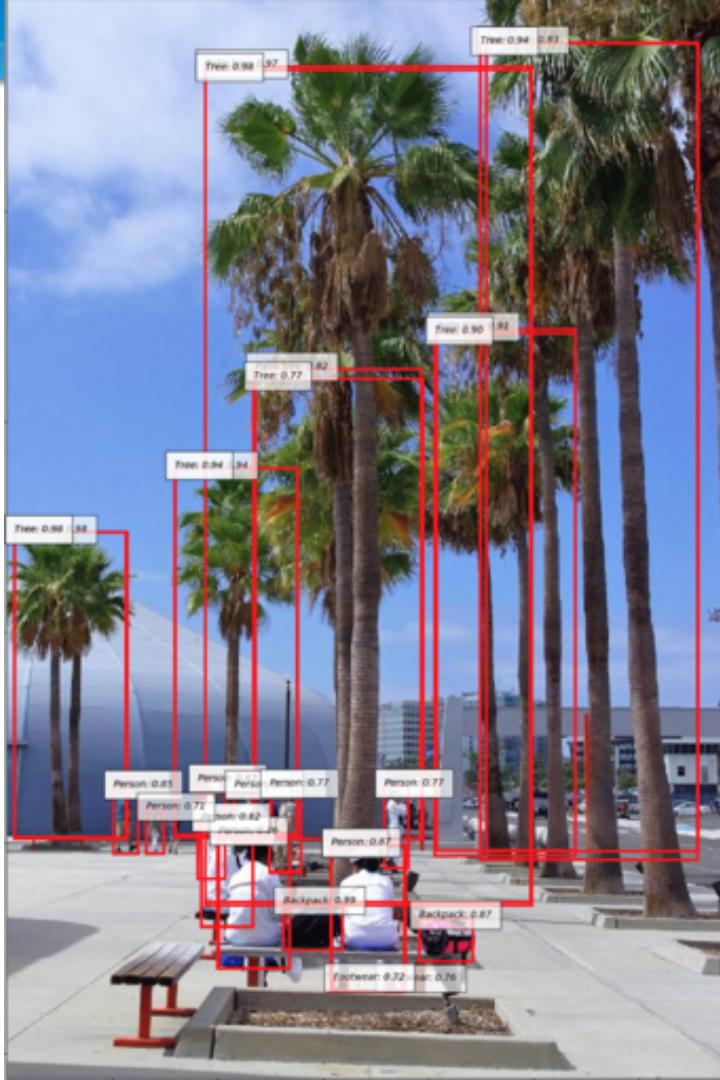


Object Detection

Detecting objects in an image
by predicting...

- bounding boxes that contain them
- category of the objects

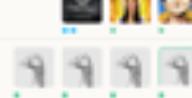




referred
networks

Achievement on MN-1b: PFDet in OIC 2018

- We won the 2nd position (0.023% diff to the 1st)

#	△pub	Team Name	Kernel	Team Members	Score	Entries	Last
1	+2	krvajnk			0.58657	102	2mo
2	-1	PFDet			0.58634	49	2mo
3	-1	Avengers			0.58616	64	2mo
4	-	XJTU			0.58348	22	2mo
5	-	ikciting			0.56801	39	2mo
6	-	Sogou_MM			0.53909	105	2mo
7	-	QLearning			0.53709	20	2mo

Congratulations to the winners again!

Object Detection track

1st place: klvajok



2d place: PFDet



3d place: Avengers



Visual Relationship Detection
track

1st place: Seiji

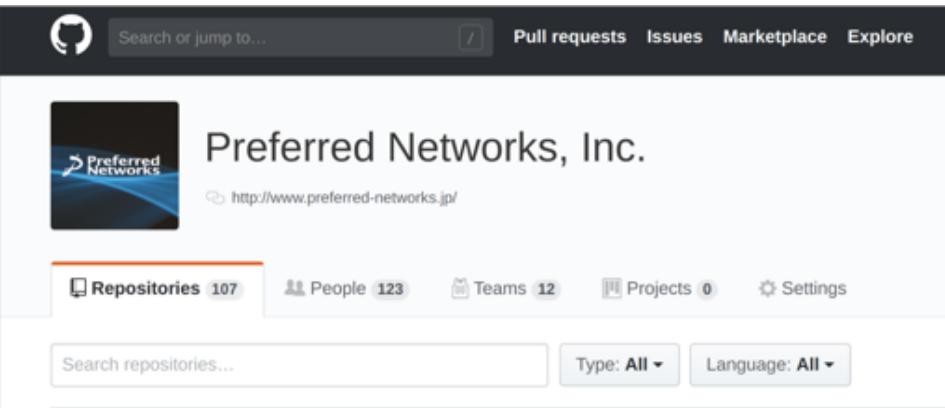


2d place: tito (individual)

3d place: Kyle (no disclosure)

4th place: toshif (individual)

Open Sourcing PFDet:



We may make the implementation public

PFDet: 2nd Place Solution to Open Images Challenge 2018
Object Detection Track

Takuya Akiba* Tommi Kerola* Yusuke Nittani* Toru Ogawa* Shotaro Sano* Shuji Suzuki*
Preferred Networks, Inc.
[\[akiba,tommi,nittani,ogawa,sano,suzuki\]@preferred.jp](mailto:{akiba,tommi,nittani,ogawa,sano,suzuki}@preferred.jp)

Abstract

We present a large-scale object detection system by team PFDet. Our system enables training with huge datasets using 512 GPUs, handles sparsely verified classes, and massive class imbalance. Using our method, we achieved 2nd place in the Google AI Open Images Object Detection Track 2018 on Kaggle.¹

1. Introduction

Open Images Detection Dataset V4 (OID) [6] is currently the largest publicly available object detection dataset, including 1.7M annotated images with 12M bounding boxes. The diversity of images in training datasets is the driving force of the generalizability of machine learning models. Successfully trained models on OID would push the frontier of object detectors with the help of data.

Training a deep learning model on OID with low parallelization would lead to prohibitively long training times, as is the case for training with other large-scale datasets [2]. We follow the work of MegNet [11] and use multi-node batch normalization to stably train an object detector with a batch size of 512. Using ChainableMN [1], a distributed deep learning library, we demonstrate highly scalable parallelization over 512 GPUs.

OID is different from its predecessors, such as MS COCO [8], not merely in terms of the sheer number of images, but also regarding the annotation style. In the previous, instances of all classes covered by the dataset are always exhaustively annotated, whereas in OID, for each image, instances of classes not verified to exist in the image are not annotated. This is a realistic approach to expanding the number of classes covered by the dataset, because without sparsifying the annotated classes, the number of annotations required may explode as the total number of classes increases.

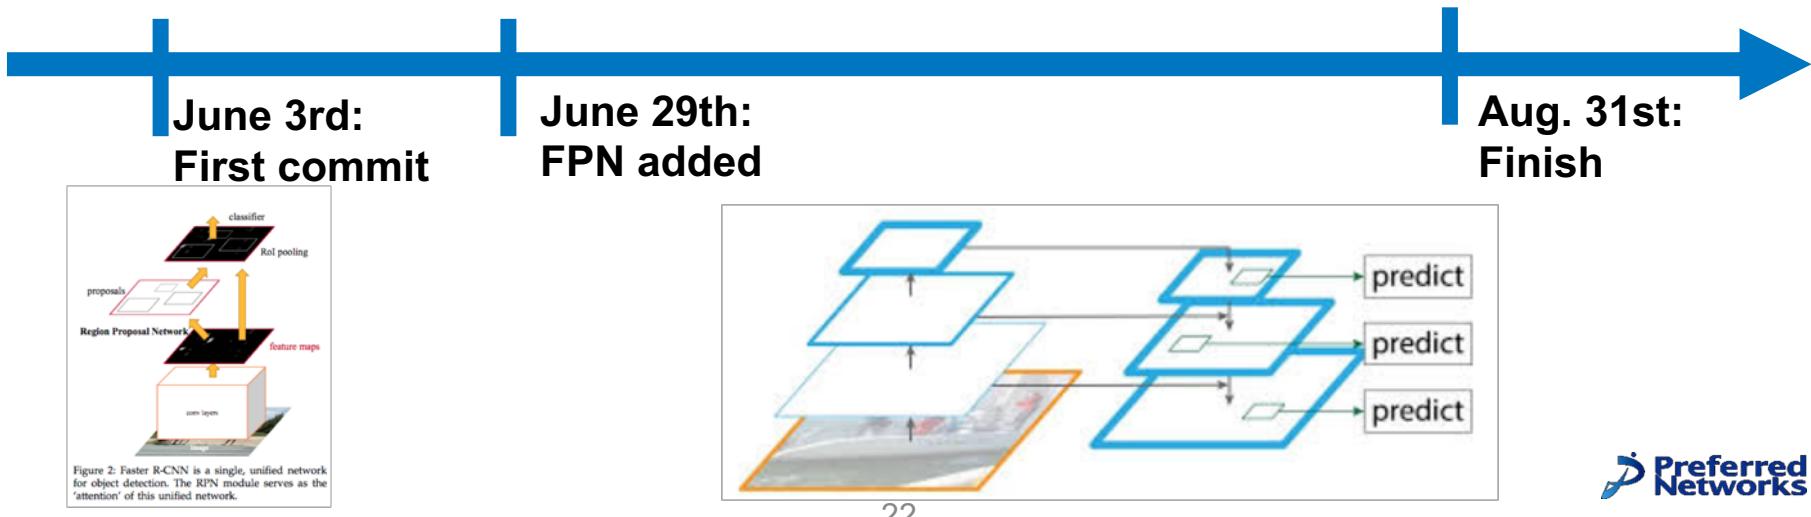
¹The authors contributed equally and they are ordered alphabetically.
²<https://www.kaggle.com/c/google-ai-open-images-object-detection-track>

Technical report is already on arXiv:
arXiv:1809.00778



Computation resources used in PFDet

- Single training process of 16 epochs takes **33 hours** with **512 x V100** GPUs of MN-1b
- Repeated model development & parameter tuning



Autonomous Tyding-up robot

Integration of a wide range of DL:

- Object Detection
(based on PFNet)
- Audio recognition
- NLP
- Picking planning

Technical topics

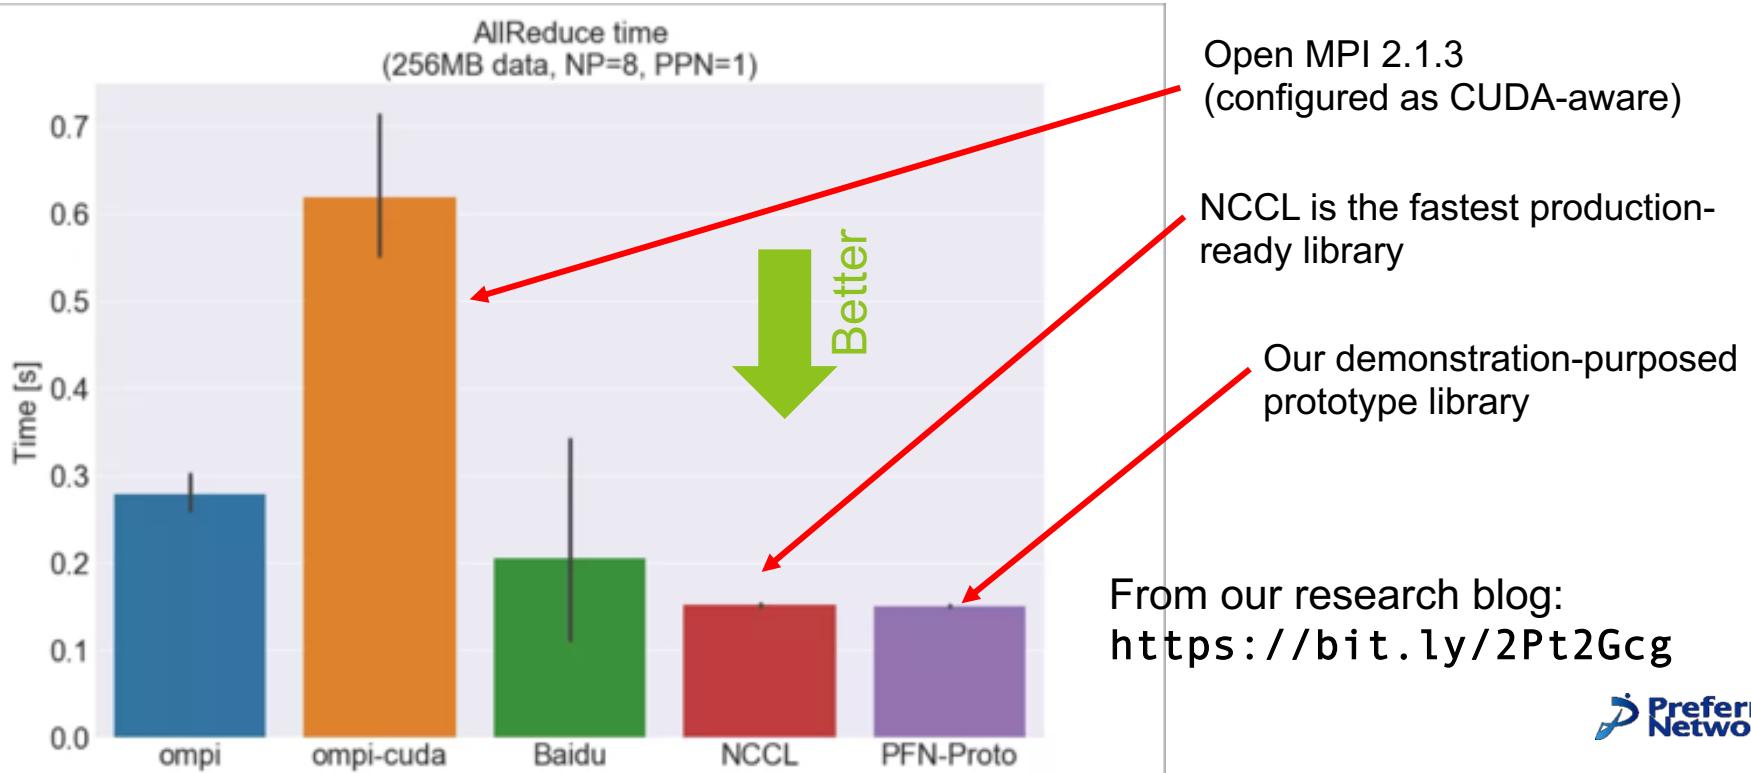
1. Communication & Fault tolerance
2. Storage
3. Resource Management

Communication

Allreduce: that always matters

- Critical component for distributed deep learning
- NVIDIA NCCL is the de-facto standard communication library
- Typical usage in ChainerMN:
 - MPI for process coordination
 - NCCL for actual communication

Communication



Communication

NCCL is now open-sourced 

- Very interesting internal architecture
- NCCL still has some bug in large scale execution... we know what to do to an OSS software? ☺

Scaling Deep Learning Training with NCCL

By Sylvain Jeaugey | September 26, 2018

Tags: Accelerated Computing, DGX-1, DGX-2, InfiniBand, NCCL, NVLink, open source

[NVIDIA Collective Communications Library \(NCCL\)](#) provides optimized implementation of inter-GPU communication. Developers using deep learning frameworks can rely on NCCL's highly optimized, MPI compatible and topology-aware communication primitives to efficiently utilize available GPUs within and across multiple nodes.

NCCL is optimized for high bandwidth and low latency over PCIe and NVLink high speed interconnect for intra-node communication and for inter-node communication. NCCL—allows CUDA applications and DL frameworks in particular—to efficiently implement complex communication algorithms and adapt them to every platform.

The latest NCCL 2.3 release makes NCCL fully open-source and available on [GitHub](#). The pre-built and tested binaries are available on [Developer Zone](#). This should provide you with the flexibility you need and enable us to have open discussions about the library.

Communication: open questions (1)

- NCCL is optimized for bandwidth-bound (i.e. large buffer) communication
- What about short-messages?
 - Inter-process Batch Normalization
 - Model parallelism
- Fault tolerance?
 - NCCL now supports communication timeout...
 - But the MPI standard does not support FT (yet?)

Communication: open questions (2)

- Do we still use MPI...?
 - Legacy
 - The software stack is huge: installation and maintenance is hard for non-HPC users
 - No fault tolerance
 - No separation of process coordination and communication
 - Hard to use with Kubernetes

Resource Management

- We deploy Kubernetes on our computing cluster.
- Why Kubernetes?
 - Advanced features from the cloud computing community
 - Fault tolerance, dynamic job size management, preemption, etc.
 - We also deploy server-based services (such as JupyterHub, CI, etc.) on the same cluster
- Still many challenges
 - Batch job scheduling,



Resource Management: Open Questions

- Preemption is a strong tool for efficient resource management
 - How DL framework and scheduler cooperate ?
 - Job preemption/restarting & flexible resource (re-)allocation
- Resource isolation
 - Bandwidth isolation for Infiniband?
 - NUMA- / topology- aware pod placement?

NOTE: We operate an in-house cluster, so all jobs are our own.

Conclusion & Takeaways:

- PFN runs a cluster of 1,500 NVIDIA GPUs
- We develop real-world solutions as well as tackle challenging competitions.
- We deploy advanced & challenging software to manage the computing resources
- And... we are hiring!

Questions?

If you have in-depth technical questions, please send it to

Keisuke Fukuda <kfukuda@preferred.jp>

Thank you!

